

Machine Learning Compendium

Machine Learning University Society at Karlsruhe Institute of
Technology (KIT)

Machine Learning Compendium

ABSTRACT

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Morbi commodo, ipsum sed pharetra gravida, orci magna rhoncus neque, id pulvinar odio lorem non turpis. Nullam sit amet enim. Suspendisse id velit vitae ligula volutpat condimentum. Aliquam erat volutpat. Sed quis velit. Nulla facilisi. Nulla libero. Vivamus pharetra posuere sapien. Nam consectetur. Sed aliquam, nunc eget euismod ullamcorper, lectus nunc ullamcorper orci, fermentum bibendum enim nibh eget ipsum. Donec porttitor ligula eu dolor. Maecenas vitae nulla consequat libero cursus venenatis. Nam magna enim, accumsan eu, blandit sed, blandit a, eros.

Contents

1	REINFORCEMENT LEARNING	4
1.1	Bellman Equations	5
1.2	Advantage Function	6
1.3	Policy, Policy Gradient	7
2	STATISTICS	8
2.1	Probability	9
2.2	Distributions	9
2.3	Estimation	10
2.4	Divergences	10
2.5	Information Geometry	12

Authors

The following authors contributed to Chapter 1: Sandro Braun

The following authors contributed to Chapter 2: Leander Kurscheidt

1

Reinforcement Learning

Reinforcement learning is an area of machine learning concerned with how software agents ought to take actions in an environment so as to maximize some notion of cumulative reward. [wikipedia](#)

1.1 BELLMAN EQUATIONS

TODO backup diagrams

1.1.1 STATE VALUE FUNCTION

$$v_{\pi}(s) = \sum_{a \in A} \pi(s|a) Q_{\pi}(s, a) \quad (1.1)$$

1.1.2 ACTION VALUE FUNCTION

$$Q_{\pi}(s, a) = r_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_{\pi}(s') \quad (1.2)$$

1.1.3 STATE VALUE FUNCTION RECURSIVE

$$v_{\pi}(s) = \sum_{a \in A} \pi(s|a) (r_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_{\pi}(s')) \quad (1.3)$$

1.1.4 ACTION VALUE FUNCTION RECURSIVE

$$Q_{\pi}(s, a) = r_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \sum_{a' \in A} \pi(a'|s') Q_{\pi}(s', a') \quad (1.4)$$

1.1.5 OPTIMAL STATE VALUE FUNCTION

$$v_*(s) = \max_a Q_*(s, a) \quad (1.5)$$

1.1.6 OPTIMAL ACTION STATE VALUE FUNCTION

$$Q_*(s, a) = r_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s') \quad (1.6)$$

1.1.7 OPTIMAL STATE VALUE FUNCTION RECURSIVE

$$v_*(s) = \max_a r_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s') \quad (1.7)$$

1.1.8 OPTIMAL ACTION STATE VALUE FUNCTION RECURSIVE

$$Q_*(a, s) = r_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \max_{a'} Q_*(s', a') \quad (1.8)$$

1.2 ADVANTAGE FUNCTION

TODO

1.3 POLICY, POLICY GRADIENT

1.3.1 POLICY: DISTRIBUTION OVER ACTIONS GIVEN STATES

$$\pi_{\theta}(a|s) = P(a|s) \tag{1.9}$$

1.3.2 POLICY GRADIENT

$$\nabla_{\theta} \pi_{\theta}(s|a) = \pi_{\theta}(s|a) \nabla_{\theta} \log \pi_{\theta}(s|a) \tag{1.10}$$

Note: this is valid for all probability distributions (the policy is a distribution over actions given states). The gradient term on the right hand side is called score function. The derivation basically uses the "log-trick".

2

Statistics

Statistics is a branch of mathematics dealing with the collection, organization, analysis, interpretation and presentation of data.

[wikipedia](#)

2.1 PROBABILITY

TODO: Probability (general + simple), CDF, Variance, Fisher-Information etc.

2.1.1 L_p -SPACE FOR RANDOM-VARIABLES

The L_p -Norm for Random-Variables X , where $\mathbb{E}|X|^p < \infty$, is defined through:

$$\|X\|_p := (\mathbb{E}[|X|^p])^{\frac{1}{p}}$$

lecture

2.1.2 JENSENS-INEQUALITY FOR RANDOM VARIABLES

If ϕ is a konvex function and X a Random-Variable, then

$$\phi(\mathbb{E}X) \leq \mathbb{E}\phi(X)$$

wikipedia

2.2 DISTRIBUTIONS

TODO: All the usual suspects, including observed

2.3 ESTIMATION

TODO: ML, Score-Function, biased/unbiased, CramérRao bound, confidence-interval

2.4 DIVERGENCES

Conventions for this section: P and Q are probability measures over a set X , and P is absolutely continuous with respect to Q . S is a space of all probability distributions with common support.

2.4.1 DIVERGENCE

A divergence on S is a function $D : S \times S \rightarrow R$ satisfying

1. $D(p||q) \geq 0 \forall p, q \in S$,
2. $D(p||q) = 0 \Leftrightarrow p = q$

A divergence is a "sense" of distance between two probability distributions. It's not a metric, but a pre-metric.

[wikipedia](#)

2.4.2 F-DIVERGENCE

1. Generalization of whole family of divergences
2. For a convex function f such that $f(1) = 0$, the f-divergence of P from Q is defined as:
$$D_f(P \parallel Q) \equiv \int_{\Omega} f\left(\frac{dP}{dQ}\right) dQ$$

3. [wikipedia](#)

2.4.3 KL-DIVERGENCE

1. The KullbackLeibler divergence from Q to P is defined as

$$D_{\text{KL}}(P\|Q) = \int_X \log \frac{dP}{dQ} dP = D_t \log t.$$

2. maximizing likelihood is equivalent to minimizing $D_{\text{KL}}(P(\cdot|\theta^*) \| P(\cdot|\theta))$, where $P(\cdot|\theta^*)$ is the true distribution and $P(\cdot|\theta)$ is our estimate.

3. [wikipedia](#)

4. TODO: Fisher-Matrix infinitesimal relationship

2.4.4 JENSEN SHANNON DIVERGENCE

The JensenShannon divergence from Q to P is defined as

$$\text{JSD}(P \| Q) = \frac{1}{2}D(P \| M) + \frac{1}{2}D(Q \| M)$$

,

where $M = \frac{1}{2}(P + Q)$

[wikipedia](#)

2.4.5 TODO: WASSERSTEIN & WASSERSTEIN DUAL

2.5 INFORMATION GEOMETRY

TODO